

# Bayesian Analysis using Stan and brms

Xiaochuan Li  
2022-06-02


# Agenda

- Bayesian Analysis
- Software
- brms
- Examples

## Frequentist vs. Bayesian analysis

- Frequentist Statistics
  - highest probability of seeing the data given (conditional on) the underlying model
- Bayesian Statistics
  - Which model, or set of parameter values, are most certain to be the true model
  - Bayes' Theorem

$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)}$$



## Why use Bayesian analysis

- Prior Information
- Hierarchical Models
- more interpretable confidence intervals
- Limited Data
  - Prior distribution over the parameters can act as a regularization to prevent unlikely extreme values

## Software

### R interface

- lme4, rethinking::quap - Quadratic approximation
- BUG, JAGS - Gibbs sampling
- STAN - Hamiltonian Sampling

### Python interface

- PyMc3 - Hamiltonian Monte Carlo
- Pytorch::Pyro - Bayesian statistical models and Neural Networks

# Stan

## Linear model

$$\log(y_i) \sim N(b_0 + b_1 \log(x_i), \sigma^2)$$

### data block

```
data {  
  int<lower=1> N; // number of observations  
  vector[N] log_gest; // log gestational age  
  vector[N] log_weight; // log birth weight  
}
```

### parameters block

```
parameters {  
  vector[2] beta; // coefs  
  real<lower=0> sigma; // error sd for Gaussian likelihood  
}
```

### model block

```
model {  
  // Log-likelihood  
  target += normal_lpdf(log_weight | beta[1] + beta[2] * log_gest, sigma);  
  // Log-priors  
  target += normal_lpdf(sigma | 0, 1)  
    + normal_lpdf(beta | 0, 1);  
}
```

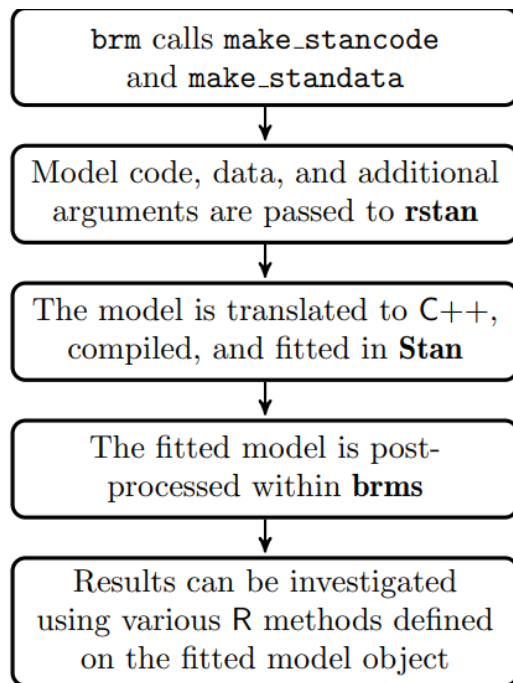
### generated quantities block

```
generated quantities {  
  vector[N] log_lik; // pointwise log-likelihood for LOO  
  vector[N] log_weight_rep; // replications from posterior predictive distribution  
  for (n in 1:N) {  
    real log_weight_hat_n = beta[1] + beta[2] * log_gest[n];  
    log_lik[n] = normal_lpdf(log_weight[n] | log_weight_hat_n, sigma);  
  }  
}
```

## brms

brms provides an interface to fit Bayesian generalized (non-)linear multivariate multilevel models using Stan.

- brms ---> Stan ---> C++



2

High level description of the model fitting procedure used in **brms**.

- Domino environment: DAD\_STAN

## brms formula

```
fit1 <- brm(formula = time | cens(censored) ~ age * sex + disease + (1 + age|patient),  
data = kidney, family = lognormal(),  
prior = c(set_prior("normal(0,5)", class = "b"),  
set_prior("cauchy(0,2)", class = "sd"),  
set_prior("lkj(2)", class = "cor")),  
warmup = 1000, iter = 2000, chains = 4,  
control = list(adapt_delta = 0.95))
```

- formula syntax matches that of lme4
- form | fun(variable)
  - cens: handles censored data
  - weights and disp: allow different sorts of weighting
  - se: specify known standard errors primarily for meta-analysis
  - trunc: define truncation boundaries
  - trials: for binomial models
  - cat: specify the number of categories for ordinal models
- Group-level terms are of the form (coefs | group)



## brms family

- Linear and robust linear regression: gaussian or student family with the identity link
- categorical - bernoulli, binomial, and categorical with logit link
- count data - poisson, negbinomial, and geometric
- survival regression - lognormal, Gamma, exponential, and weibull
- Ordinal regression - cumulative, cratio, sratio, and acat
- Excess zeros response - zero\_inflated\_poisson, zero\_inflated\_negbinomial, zero\_inflated\_binomial, zero\_inflated\_beta, hurdle\_poisson, hurdle\_negbinomial, and hurdle\_gamma
- Specify custom families to include more distributions

## brms parameters

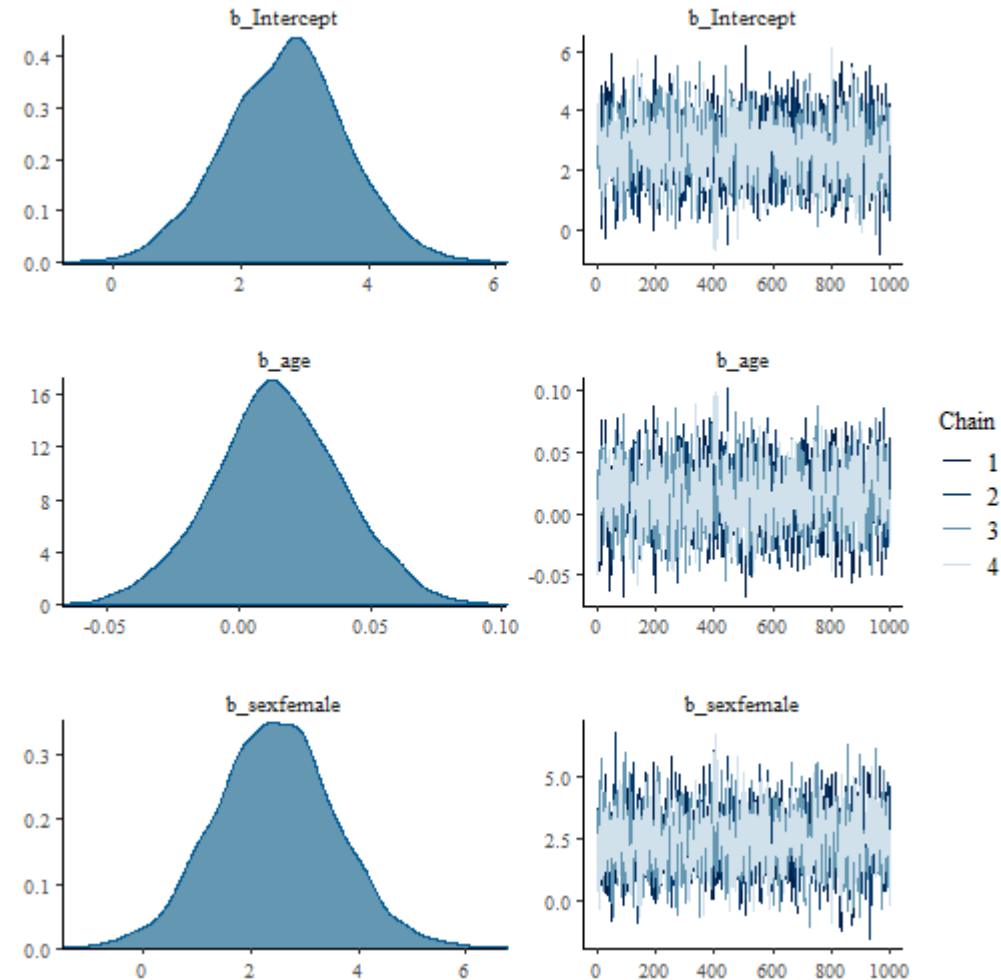
- prior: Prior distributions of model parameters

```
prior <- c(set_prior("normal(0,10)", class = "b", coef = "age"),  
set_prior("cauchy(1,2)", class = "b", coef = "sexfemale"),  
set_prior("student_t(3,0,5)", class = "sd", group="patient"))
```

- control: Adjusting the sampling behavior of Stan
  - choosing the number of iterations, warmup samples, and chains,
  - decrease (or eliminate at best) the number of divergent transitions that cause a bias in the obtained posterior samples. "There were x divergent transitions after warmup."

## Check Posterior distribution

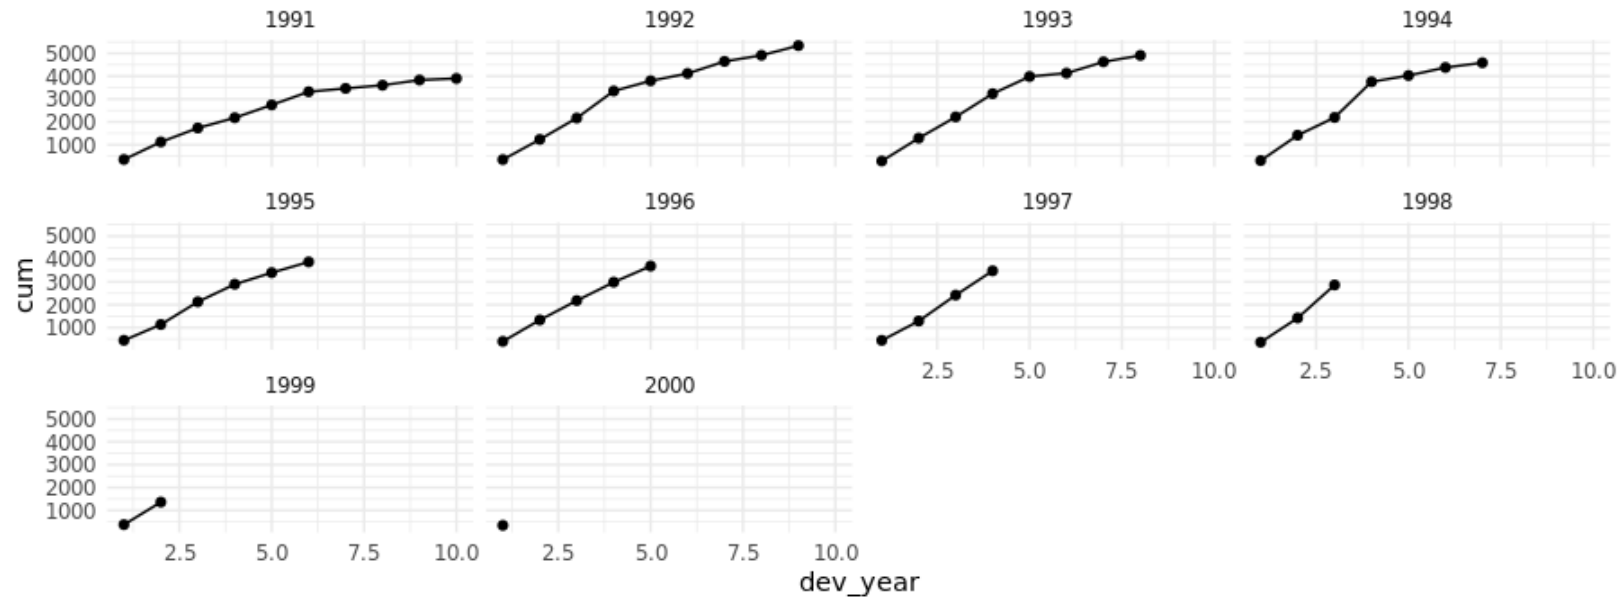
Posterior distribution of parameters and chain convergence need to be checked



# Example

## non-linear model to model insurance loss

- Data



- Model

$$cum_{AY,dev} \sim N(\mu_{AY,dev}, \sigma)$$
$$\mu_{AY,dev} = ult_{AY} \left( 1 - \exp\left(-\left(\frac{dev}{\rho}\right)^w\right) \right)$$

## Prior parameter distributions

- Prior parameter distribution

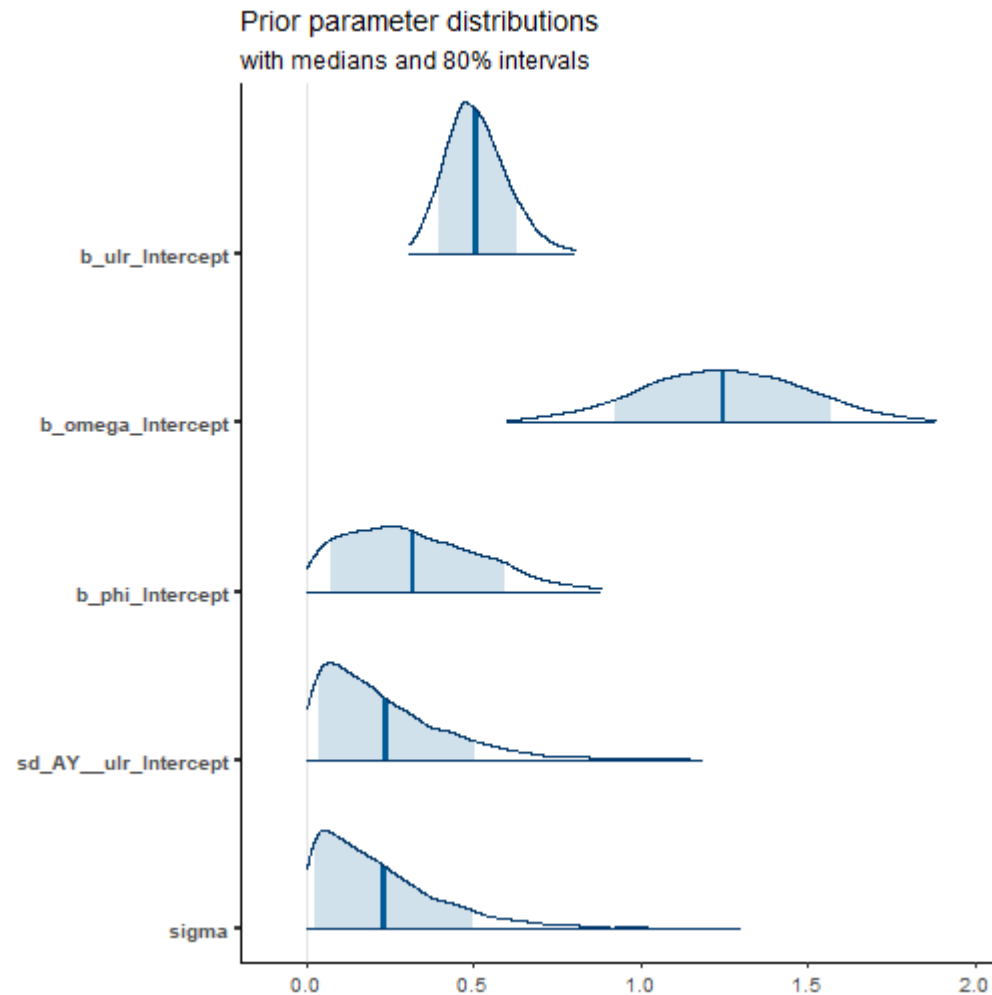
$$\phi \sim N(0.25, 0.25^2)^+$$

$$\omega \sim N(1.25, 0.25^2)^+$$

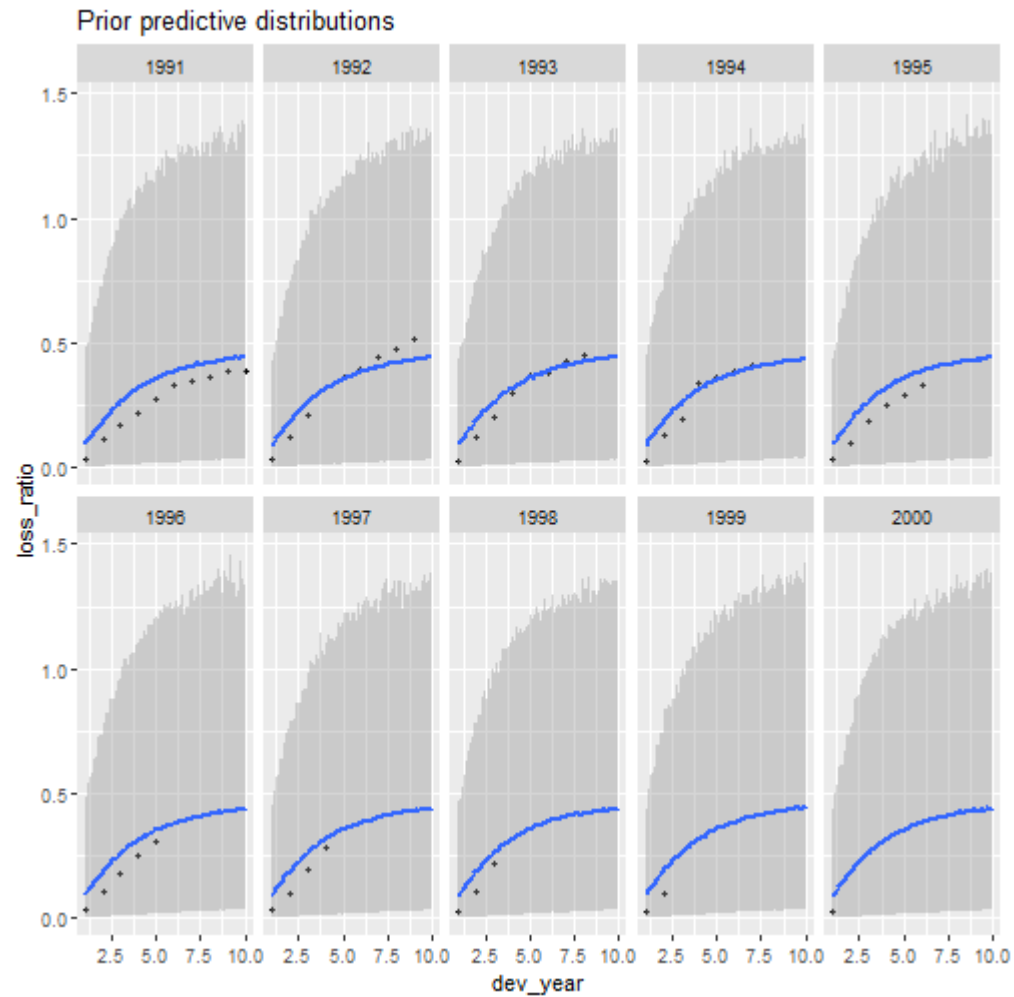
$$\gamma \sim \text{logN}(\text{log}(0.5), \text{log}(1.2)^2)$$

$$\sigma \sim \text{Student-t}(5, 0, 0.25)^+$$

$$\sigma_{\gamma^0} \sim \text{Student-t}(5, 0, 0.25)^+$$



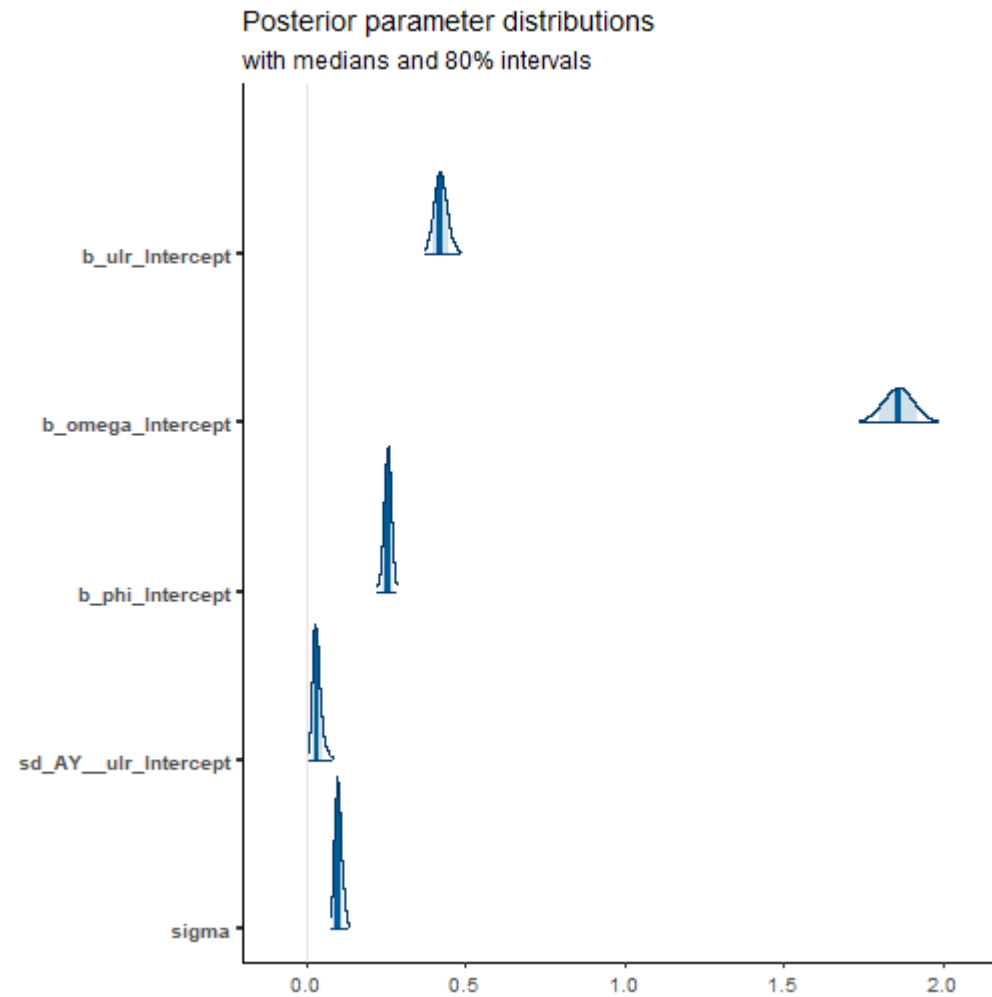
## Estimate on Prior predictive distribution



# Model

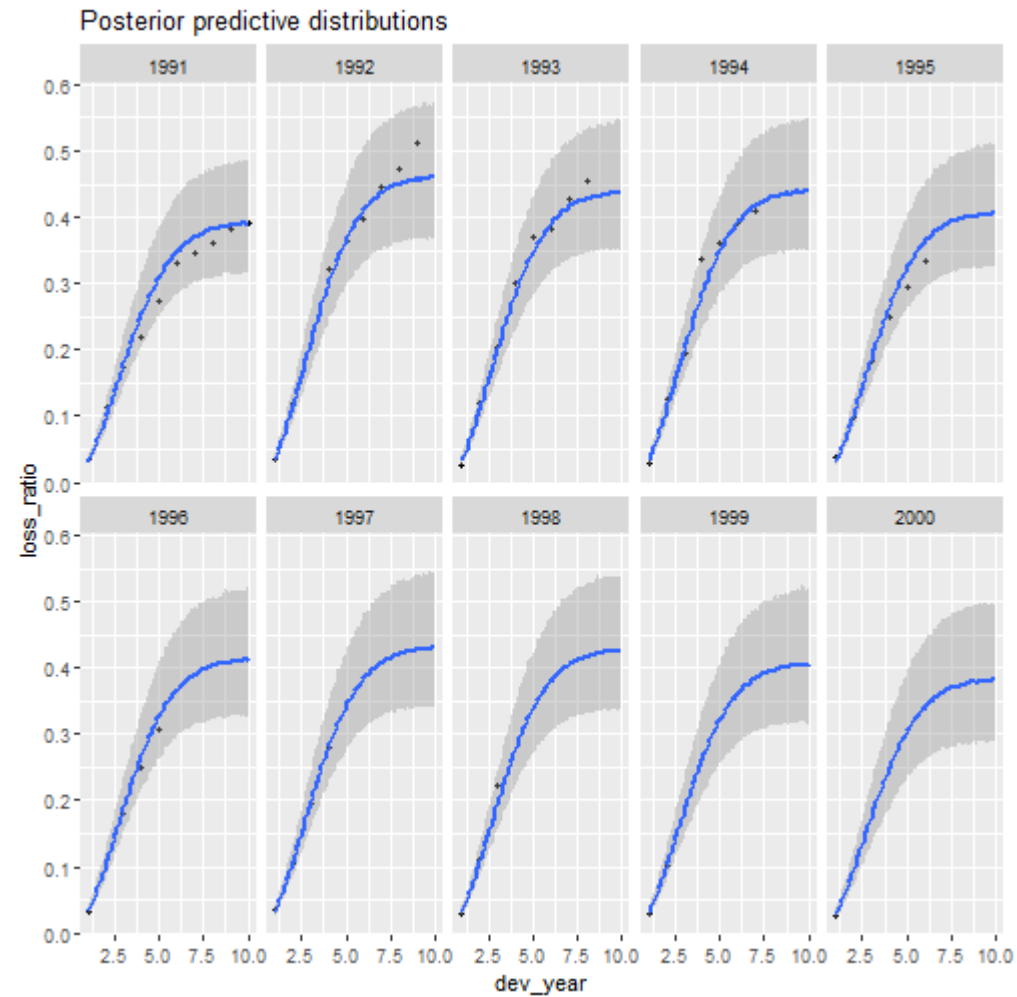
```
nlform2 <- bf(loss_ratio ~ log(ulr * (1 - exp(-(dev_year*phi)^omega))),  
ulr ~ 1 + (1|AY), omega ~ 1, phi ~ 1,  
nl = TRUE)  
  
m2 <- brm(nlform2, data = loss,  
family = lognormal(link = "identity", link_sigma = "log"),  
prior = c(  
prior(lognormal(log(0.5), log(1.2)), nlpar = "ulr", lb=0),  
prior(normal(1.25, 0.25), nlpar = "omega", lb=0),  
prior(normal(0.25, 0.25), nlpar = "phi", lb=0),  
prior(student_t(5, 0, 0.25), class = "sigma"),  
prior(student_t(5, 0, 0.25), class = "sd", nlpar="ulr")  
),  
sample_prior = "no", seed = 1234  
)
```

## Posterior parameter distributions





## Posterior predictive distributions



# Pros and Cons of Bayesian method

- Pros
  - more thoughtful on data generating models/ process
  - Include prior knowledge
  - more flexible in model structure
- Cons
  - different mindset and could be difficult to set up
  - more time consuming to fit models
  - hard to scale to very large data

# links

<https://mc-stan.org/>

<https://cran.r-project.org/web/packages/brms/index.html>